



# LoRA: Low-Rank Adaptation of Large Language Models

Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu,  
Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen

Microsoft, ICLR 2022

# Introduction

## Pre-trained and Fine-tuning

Many applications in natural language processing rely on applying a large-scale, pre-trained language model to multiple downstream applications through fine-tuning.

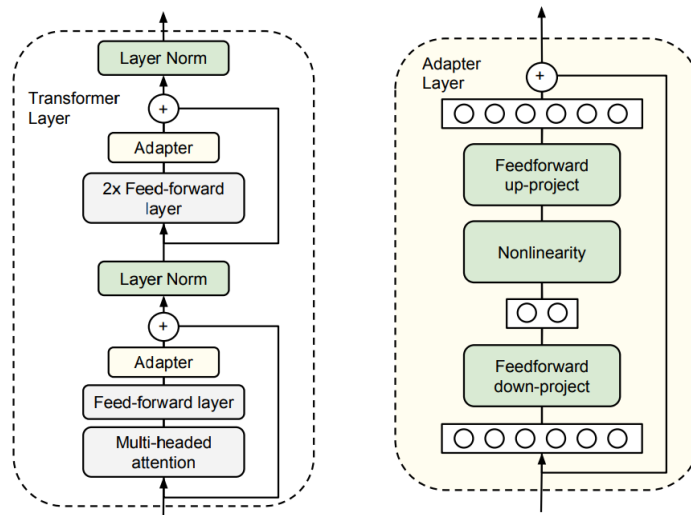
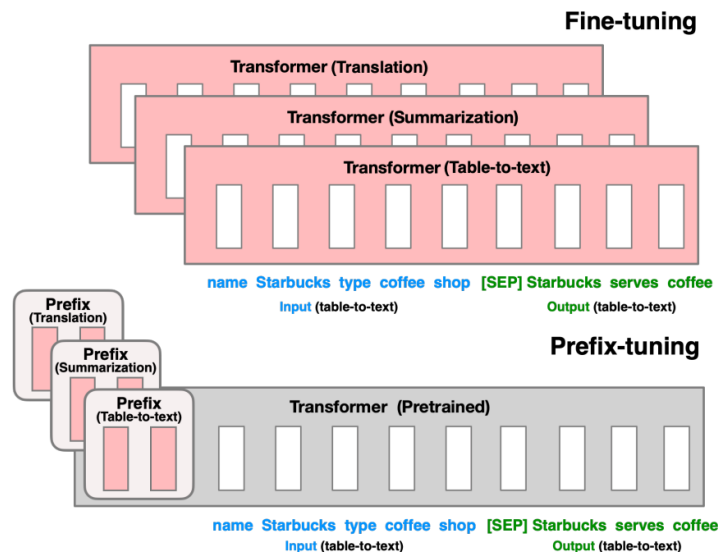
However, the fine-tuned model according to different tasks has the same size parameters as the pre-trained model, This requires a lot of space to store the models of these specialized tasks.

# Introduction (cont.)

## Existing Solution

Many sought to mitigate this by adapting only some parameters or learning external modules for new tasks.

e.g., Bias-only Tuning, Prefix-tuning, Adapter Layer.




# Introduction (cont.)

## Aren't Existing Solutions Good Enough ?

Nevertheless, techniques prior to LoRA often introduce

- inference latency by extending model depth,
- reduce the model's usable sequence length.

more importantly, these methods often fail to match the fine-tuning baselines, posing a trade-off between efficiency and model quality.




When the generation length is shorter and the batch is smaller, the additional computational cost of the Adapter Layer is more obvious.

# Low Rank Adaptation

Write the weight matrix of fine-tuning in the form of  $W + \Delta W$ , and  $W, \Delta W \in \mathbb{R}^{d \times k}$ ;  $W$  is the pre-trained weight matrix.

LoRA constrain fine-tuned update  $\Delta W$  by representing the latter with a low-rank decomposition  $\Delta W = \frac{\alpha}{r} B A^\top$ ,  $A \in \mathbb{R}^{k \times r}$ ,  $B \in \mathbb{R}^{d \times r}$  and  $r \ll \min(d, k)$ ,  $\alpha$  is a constant scalar.



# Low Rank Adaptation (cont.)

$$\Delta W \text{ of LoRA} = BA^\top$$

- LoRA's trainable parameters are much smaller than full fine-tuning due to  $r \ll d$ .

- A Generalization of Full Fine-tuning.

Roughly recover the expressiveness of full fine-tuning by setting the LoRA rank  $r$  to the rank of the pre-trained weight matrices.

- No Additional Inference Latency.

Adding the pre-training weights  $W$  to the LoRA parameters  $BA^\top$  first will not increase the amount of calculation during inference.

# Experiment

## Comparing the performance on GPT-3

Model&Method	# Trainable Parameters	WikiSQL	MNLI-m	SAMSum
		Acc. (%)	Acc. (%)	R1/R2/RL
GPT-3 (FT)	175,255.8M	<b>73.8</b>	89.5	52.0/28.0/44.5
GPT-3 (BitFit)	14.2M	71.3	91.0	51.3/27.4/43.5
GPT-3 (PreEmbed)	3.2M	63.1	88.6	48.3/24.2/40.5
GPT-3 (PreLayer)	20.2M	70.1	89.5	50.8/27.3/43.5
GPT-3 (Adapter <sup>H</sup> )	7.1M	71.9	89.8	53.0/28.9/44.8
GPT-3 (Adapter <sup>H</sup> )	40.1M	73.2	<b>91.5</b>	53.2/29.0/45.1
GPT-3 (LoRA)	4.7M	73.4	<b>91.7</b>	<b>53.8/29.8/45.9</b>
GPT-3 (LoRA)	37.7M	<b>74.0</b>	<b>91.6</b>	53.4/29.2/45.1

In the experiments of this paper, only the attention layer adds LoRA parameters.

# the Optimal Rank $r$ for LoRA

## GPT-3



	# of Trainable Parameters = 18M						
Weight Type Rank $r$	$W_q$ 8	$W_k$ 8	$W_v$ 8	$W_o$ 8	$W_q, W_k$ 4	$W_q, W_v$ 4	$W_q, W_k, W_v, W_o$ 2
WikiSQL ( $\pm 0.5\%$ )	70.4	70.0	73.0	73.2	71.4	<b>73.7</b>	<b>73.7</b>
MultiNLI ( $\pm 0.1\%$ )	91.0	90.8	91.0	91.3	91.3	91.3	<b>91.7</b>

## GPT-2

	Weight Type	$r = 1$	$r = 2$	$r = 4$	$r = 8$	$r = 64$
WikiSQL( $\pm 0.5\%$ )	$W_q$	68.8	69.6	70.5	70.4	70.0
	$W_q, W_v$	73.4	73.3	73.7	73.8	73.5
	$W_q, W_k, W_v, W_o$	74.1	73.7	74.0	74.0	73.9
MultiNLI ( $\pm 0.1\%$ )	$W_q$	90.7	90.9	91.1	90.7	90.7
	$W_q, W_v$	91.3	91.4	91.3	91.6	91.4
	$W_q, W_k, W_v, W_o$	91.2	91.7	91.7	91.5	91.4



## Experiment (cont.)



- Prefix-tuning is difficult to optimize and that its performance changes non-monotonically in trainable parameters.
- LoRA has stable performance under various number of parameters.

# Experiment (cont.)

## Low-Data Regime

Method	MNLI(m)-100	MNLI(m)-1k	MNLI(m)-10k	MNLI(m)-392K
GPT-3 (Fine-Tune)	60.2	<b>85.8</b>	88.9	89.5
GPT-3 (PrefixEmbed)	37.6	75.2	79.5	88.6
GPT-3 (PrefixLayer)	48.3	82.5	85.9	89.6
GPT-3 (LoRA)	<b>63.8</b>	85.6	<b>89.2</b>	<b>91.7</b>

Compared with the Prefix-tuning method with obvious performance degradation, LoRA still shows good performance with only a small amount of data.

# Experiment (cont.)

## Combined with Other Adaptation Methods

Method	Hyperparameters	# Trainable Parameters	WikiSQL	MNLI-m
Fine-Tune	-	175B	73.8	89.5
LoRA	$r_v = 2$	4.7 M	73.4	<b>91.7</b>
	$r_q = r_v = 1$	4.7 M	73.4	91.3
	$r_q = r_v = 2$	9.4 M	73.3	91.4
	$r_q = r_k = r_v = r_o = 1$	9.4 M	74.1	91.2
	$r_q = r_v = 4$	18.8 M	73.7	91.3
	$r_q = r_k = r_v = r_o = 2$	18.8 M	73.7	<b>91.7</b>
	$r_q = r_v = 8$	37.7 M	73.8	<b>91.6</b>
	$r_q = r_k = r_v = r_o = 4$	37.7 M	74.0	<b>91.7</b>
	$r_q = r_v = 64$	301.9 M	73.6	91.4
LoRA+PE	$r_q = r_v = 8, l_p = 8, l_i = 4$	37.8 M	75.0	91.4
	$r_q = r_v = 32, l_p = 8, l_i = 4$	151.1 M	<b>75.9</b>	91.1
	$r_q = r_v = 64, l_p = 8, l_i = 4$	302.1 M	<b>76.2</b>	91.3

# Conclusions

This study proposes LoRA to replace fine-tuning

- With the same pre-trained weights, only a small number of LoRA parameters need to be saved for different tasks.
- That neither introduces inference latency nor reduces input sequence length while retaining high model quality.
- Can be combined with other efficient adaptation methods, potentially providing orthogonal improvement.

## Next Version

AdaLoRA, which adaptively allocates the parameter budget among weight matrices according to their importance score.